

Reference Manual for the Java Modeling Language

David R. Cok, Gary Leavens, Mathias Ulbrich, TBD

DRAFT December 21, 2016

Copyright (c) 2010-2017 TBD

Contents

1	Introduction	2
1.1	Behavioral Interface Specifications	2
1.2	A First Example	2
1.3	What is JML Good For?	2
1.4	Status, Plans and Tools for JML	2
1.5	Previous JML reference manual	2
1.6	Historical Precedents and Antecedents	3
1.7	Acknowledgments	3
2	Structure of this Manual	4
2.1	Typographical conventions	4
2.2	Grammar	4
3	JML concepts	6
3.1	JML and Java compilation units	6
3.2	JML specifications and Java files	7
3.3	.jml files	8
3.3.1	Combining Java and JML files	9
3.4	Specification inheritance	10
3.5	JML modifiers and Java annotations	11
3.5.1	Modifiers	11
3.5.2	Type modifiers	12
3.6	Model and Ghost	12
3.7	Visibility	12
3.8	Evaluation and well-formedness of JML expressions	12
3.9	Null and non-null references	12
3.10	Static and Instance	12
3.11	Observable purity	13
3.12	Location sets and Dynamic Frames	13
3.13	Arithmetic modes	13
3.14	Immutable types and functions	13
3.15	Race condition detection	13
3.16	Redundant specifications	13
3.17	Controlling warnings	13

3.18	org.jmlspecs.lang package	13
3.19	Interaction with other tools	14
3.19.1	Interaction with Type Annotations in Java 1.8	14
3.19.2	Interaction with the Checker framework	14
3.19.3	Interaction with FindBugs	14
4	JML Syntax	15
4.1	Textual form of JML specifications	15
4.2	JML Syntax	16
4.2.1	Syntax of JML specifications	16
4.2.2	Conditional JML specifications	17
5	JML Types	19
5.1	\bigint	19
5.2	\real	19
5.3	\TYPE	20
5.3.1	JML types	20
6	JML Specifications at Package and File level	21
6.1	Model import statements	21
6.2	Default imports	22
6.3	Issues with model import statements	22
7	Specifications of Java types	23
7.1	non_null_by_default, nullable_by_default, @NonNullByDefault, @NullableByDefault	23
7.2	invariant clause	23
7.3	constraint clause	23
7.4	initially clause	24
7.5	ghost fields	24
7.6	model fields	24
7.7	represents clause	24
7.8	model methods and model classes	24
7.9	initializer and static_initializer	24
7.10	axiom	24
7.11	readable if clause and writable if clause	24
7.12	monitors_for clause	25
8	JML Method specifications	26
8.1	Structure of JML method specifications	26
8.1.1	Behaviors	27
8.1.2	Nested specification clauses	28
8.1.3	Specification Inheritance	28
8.1.4	Modifiers and annotations	29
8.1.5	Visibility	29
8.1.6	Grammar of method specifications	30

8.2	Method specifications as Annotations	30
8.3	Modifiers for methods	30
8.4	Common JML method specification clauses	30
8.4.1	requires clause	30
8.4.2	ensures clause	30
8.4.3	assignable clause	31
8.4.4	signals clause	31
8.4.5	signals_only clause	31
8.5	Advanced JML method specification clauses	31
8.5.1	accessible clause	31
8.5.2	diverges clause	31
8.5.3	measured_by clause	31
8.5.4	when clause	31
8.5.5	old clause	31
8.5.6	forall clause	32
8.5.7	duration clause	32
8.5.8	working_space clause	32
8.5.9	callable clause	32
8.5.10	captures clause	32
8.6	Model Programs (model_program clause)	32
8.6.1	Structure and purpose of model programs	32
8.6.2	extract clause	32
8.6.3	choose clause	32
8.6.4	choose_if clause	32
8.6.5	or clause	33
8.6.6	returns clause	33
8.6.7	continues clause	33
8.6.8	breaks clause	33
8.7	TODO Somewhere	33
9	Field Specifications	34
9.1	Field and Variable Modifiers	34
9.1.1	non_null and nullable (@NonNull, @Nullable)	34
9.1.2	spec_public and spec_protected (@SpecPublic, @SpecProtected)	34
9.1.3	ghost and @Ghost	36
9.1.4	model and @Model	36
9.1.5	uninitialized and @Uninitialized	36
9.1.6	instance and @Instance	36
9.1.7	monitored and @Monitored	36
9.1.8	query, secret and @Query, @Secret	36
9.1.9	peer, rep, readonly (@Peer, @Rep, @Readonly)	36
9.2	Datagroups: in and maps clauses	36
9.3	Ghost fields	36
9.4	Model fields	36
10	JML Statement Specifications	37

10.1	assert statement and Java assert statement	37
10.2	assume statement	38
10.3	ghost and model declarations	38
10.4	unreachable statement	38
10.5	reachable statement	39
10.6	set and debug statements	39
10.7	statement (block) specification	39
10.8	loop specifications	40
11	JML Expressions	41
11.1	org.jmlspecs.lang.JML	41
11.2	Java operations used in JML	41
11.3	Precedence of infix operations	42
11.4	Well-defined expressions	42
11.5	Implies and reverse implies: \Rightarrow \Leftarrow	44
11.6	Equivalence and inequivalence: \Leftrightarrow \nLeftrightarrow	44
11.7	JML subtype: $<$	45
11.8	Lock ordering:	45
11.9	\result	45
11.10	\exception	46
11.11	\old, \pre, and \past	46
11.12	\lblpos, \lblneg, \lbl, and JML.lblpos(), JML.lblneg(), JML.lbl() 46	
11.13	\nonnull elements	47
11.14	\fresh	48
11.15	informal expression: $(*...*)$ and JML.informal()	49
11.16	\type	49
11.17	\typeof	50
11.18	\elemtype	50
11.19	\is_initialized	51
11.20	\invariant_for	51
11.21	\not_modified	52
11.22	\lockset and \max	52
11.23	\reach	53
11.24	\duration	53
11.25	\space	53
11.26	\working_space	53
12	Arithmetic modes	54
13	Universe types	55
14	Model Programs	56
15	Obsolete and Deprecated Syntax	57
16	Grammar Summary	58

17 Type Checking Summary	59
18 Verification Logic Summary	60
19 Differences in JML among tools	61
20 Misc stuff to move, incorporate or delete	62
20.0.1 Finding specification files and the refine statement	62
20.0.2 Model import statements	63
20.0.3 Modifiers	63
20.0.4 JML expressions	65
20.0.5 Code contracts	65
20.1 JML modifiers and Java annotations	65
20.2 org.jmlspecs.lang package	67
20.3 JML Syntax	68
20.3.1 Syntax of JML specifications	68
20.3.2 Conditional JML specifications	68
20.3.3 Finding specification files and the refine statement	69
20.3.4 Model import statements	70
20.3.5 Modifiers	71
20.3.6 JML expressions	72
20.3.7 JML types	72
21 Statement translations	73
21.1 While loop	73
22 Java expression translations	75
22.1 Implicit or explicit arithmetic conversions	75
22.2 Arithmetic expressions	75
22.3 Bit-shift expressions	76
22.4 Relational expressions	76
22.5 Logical expressions	76

General notes on things not to forget:

- enum types
- default specs for binary classes
- datagroups, JML.* utility functions, @Requires-style annotations. arithmetic modes, universe types
- interaction with JSR 308
- various @NonNull annotations in different packages
- visibility in JML
- Sorted First-order-logic
- individual subexpressions; optional expression form; optimization; usefulness for tracing
- RAC vs. ESC
- nomenclature
- lambda expressions
- other Java 6,7,8 features
- Specification of subtypes - cf Clyde Ruby's thesis

Chapter 1

Introduction

1.1 Behavioral Interface Specifications

To write - take from DRM?

1.2 A First Example

To write - take from DRM?

1.3 What is JML Good For?

To write - take from DRM?

1.4 Status, Plans and Tools for JML

To write - take from DRM? Include somewhere a discussion of available tools and other resources, to include at least OpenJML, Key, OpenJML tutorial, Key book, etc.

1.5 Previous JML reference manual

This reference manual builds on the previous draft JML Reference Manual [1], which has been evolving over many years and has many contributors. This current edition of

the reference manual is largely a rewrite of the previous draft; it incorporates the experience of building tools for JML by the OpenJML and Key developers, many decisions about new features or deprecated features made at JML workshops, and discussions about JML on the JML mailing lists. This edition of the reference manual includes features that are proposed enhancements or clarifications of the consensus language definition. It also includes rationale for non-obvious language features and discussion of points that are under current debate or require extended explanation.

JML changes as the current version of Java changes. As this document is being written, Java 8 is the current version of Java, with Java 9 on the horizon. The version of JML presented here corresponds to Java 8. Java 8 has affected JML by, for example, the introduction of type annotations and lambda expressions.

1.6 Historical Precedents and Antecedents

To write - take from DRM? Include description of other languages that have been inspired by JML for other source languages

1.7 Acknowledgments

To write - take from DRM?

Chapter 2

Structure of this Manual

[Describe overall organization](#)

2.1 Typographical conventions

The remaining chapters of this book follow some common typographical conventions.

This style of text is used for commentary on the JML language itself, such as outstanding issues or now-obsolete practice.

[comments on how grammars are written; reference to Java grammar](#)

2.2 Grammar

The grammar of JML is intertwined with that of Java. The grammar is given in this Reference Manual as extensions of the Java grammar, using conventional BNF-style productions. The meta-symbols of the grammar are in slightly larger, normal-weight, mono-spaced font. The productions of the grammar use the following syntax:

- non-terminals are written in italics and enclosed in angle brackets: *<expression>*
- terminals are written in bold typewriter font: **forall**
- parentheses express grouping: ()
- an infix vertical bar expresses mutually-exclusive alternatives: ... | ... | ...
- optional elements and repetitions of 0 or more and 1 or more use post-fixed symbols: ? * +

- a post-fixed ... indicates a comma-separated list of 0 or more elements:
 <expression> ...
- a production has the form: <non-terminal> ::= ...

Chapter 3

JML concepts

This chapter describes some general design principles and concepts of the Java Modeling Language. [Say more by way of introduction and motivation](#)

JML specifications are declarative statements about the behavior of Java entities, namely, packages, classes, methods, and fields. Typically JML does not make assertions about how a method or class is implemented, only about the net behavior of the implementation. However, to aid in proving assertions about the behavior of methods, JML does include statement and loop specifications.

Specifications and corresponding proof obligations may be considered at different levels of granularities. JML typically is concerned with modular proofs at the level of Java methods. That is, a verification system will establish that each Java method of a program is consistent with its own specifications, presuming the specifications of all methods and classes it uses are correct. If this statement is true for all methods in the program, and all methods terminate, then the system as a whole is consistent with its specifications. [Reference for this claim?](#)

[say more about what it would mean to be modular at a different granularity](#)

3.1 JML and Java compilation units

The Java programming language is organized as a set of *compilation units* grouped into packages. The Java language specification does not stipulate a particular means of storing the Java program text that constitutes each compilation unit. However, the vast majority of systems supporting Java programs store each compilation unit as a separate file; the files constituting a package are organized into directory hierarchies corresponding the package and class names.

Specifications written in JML may either (a) be part of the same text that constitutes a Java compilation unit or (b) be in a separate body of text associated with the Java

compilation unit. For Java language systems in which Java source material is stored in files, then the JML specifications are either in the same file (case (a)) or in a separate file (case (b)). In case (b), the separate file has a `.jml` suffix, the same root name as the corresponding Java source file (typically the name of the public class in the compilation unit), the same package designation, and is stored in the file system's directory hierarchy according to its package and class name, similarly to the Java compilation unit source files.

[Review the above and the next two sections for consistency, lack of overlap, etc. – they are combined from different sources.](#)

3.2 JML specifications and Java files

The simplest way of specifying a Java program with JML is to include the text of the JML specifications directly in the Java source file.

3.3 .jml files

Do we need to say anything about how .jml files are stored in the file system - according to package hierarchy, like Java files?

JML files have a .jml extension and have a similar appearance to the corresponding .java file. The form follows the following rules. Every .jml file has a corresponding .java or .class file; where no .java file is available, the rules below refer to the .java file that would have been compiled to produce the .class file.

The principle present throughout these rules is that declarations in a JML file either (1) correspond to a declaration in the Java file, having the same name, types, non-JML modifiers and annotations, or (2) do not correspond to a Java declaration, and then must have a different name. Declarations that correspond must not be in JML comments and must not be marked `ghost` or `model`; JML declarations that do not correspond to Java declarations must be in JML comments and must be marked `ghost` or `model`.

File-level rules

- The .jml file has the same package declaration as the .java file.
- The .jml file may have a different set of import statements and may, in addition, include model import statements.
- The .jml file must include a JML declaration of the public type (i.e., class or interface) declared in the .java file. It may but need not have JML declarations of non-public types present in the .java class. Any type declared in the .jml file that is not present in the .java file must be in a JML comment and must have a `model` modifier.

Class declarations

- The JML declaration of a class and the corresponding Java declaration must extend the same superclass, implement the same set of interfaces and have the same set of non-JML modifiers (`public`, `protected`, `private`, `static`, `final`, [What others](#)). The JML declaration may add additional JML modifiers or annotations.
- Nested and inner class declarations within an enclosing non-model JML class declaration must follow the same rules as file-level class declarations: they must either correspond in name and properties to a corresponding nested or inner Java class declaration or be a model class.
- JML model classes need not have full implementations, as if they were Java declarations. However, if runtime-assertion checking tools are expected to check or use the model class, it must have a compilable and executable declaration.

Interface declarations

- The JML declaration of an interface and the corresponding Java declaration must extend the same set of interfaces and have the same set of non-JML modifiers (`public`, `protected`, `private`, `static`, [What others](#)). The JML declaration may add additional JML modifiers or annotations.
- [Comment on static and instance; no initializer for JML field declarations](#)

Method declarations

- Methods declared in a non-model JML type declaration must either correspond precisely to a method declared in the corresponding Java type declaration or be a model method. *Correspond precisely* means having the same name, exactly the same argument and return types, and the same set of exceptions.
- Methods that correspond to Java methods must not be declared model and must not have a body. They must have the same set of non-JML modifiers and annotations as the Java declaration, but may add additional JML modifiers and annotations.

Field declarations

- Fields declared in a non-model JML type declaration must either correspond precisely to a field declared in the corresponding Java type declaration or be a model or ghost field. *Correspond precisely* means having the same name and type and non-JML modifiers and annotations. The JML declaration may add additional JML modifiers and annotations.
- A JML field declaration that corresponds to a Java field declaration may not be in a JML comment, may not be `model` or `ghost` and must not have an initializer.
- A JML field declaration that does not correspond to a Java field declaration must be in a JML comment and must be either `ghost` or `model`.
- `ghost` field declarations have the same grammatical form as Java declarations, except that they may use JML types and operators and may refer to names declared in other `ghost` or `model` declarations.
- `model` field declarations have the same grammatical form as Java declarations, except that they may use JML types and operators; they may not have initializers.

Initializer declarations [Say something](#)

3.3.1 Combining Java and JML files

The specifications for the Java declarations within a Java compilation unit are determined as follows. IF there is no corresponding `.jml` file, then the specifications are those present in the `.java` file; if there is no `.java` file and no `.jml` file, only a `.class` file, then default specifications are used. `jmlxtoD`Where described

If there is a `.jml` file corresponding to a `.java` file, then the JML specification present in the `.jml` file supersedes all of the JML specifications in the `.java` file; those in the `.java` file are ignored.

If there is a corresponding `.jml` and `.java` file, processing proceeds as follows. First all matches among type declarations are established recursively:

- Top-level types in each file are matched by name. The type-checking pass checks that the modifiers, superclass and super interfaces match. JML classes that match are not model and are not in JML comments; JML classes that do not match must be model and must be in JML comments. Not all Java declarations need have a match in JML; those that have no match will have default specifications. The specifications for a Java
- Model types contain their own specifications and are not subject to further matching.
- For each non-model type, matches are established for the nested and inner type declarations in the `.jml` and `.java` declarations by the same process, recursively.

Then for each pair of matching JML and Java declarations, matches are established for method and field declarations.

- Field declarations are matched by name. Type-checking assures that declarations with the same name have the same type, modifiers and annotations.
- Method declarations are matched by name and signature. This requires that all the processing of import statements and type declarations is complete so that type names can be properly resolved.

For each pair of matching declarations, the JML specifications present in the `.jml` file give the specifications for the Java entity being declared. If there is a `.jml` file but no match for a particular Java declaration in the corresponding `.java` file, then that declaration uses default specifications, even if the `.java` file contains specifications. The contents of the `.jml` file supersede all the JML contents of the `.java` file; there is no merging of the files' contents.¹

3.4 Specification inheritance

Object-oriented programming with inheritance requires that derived classes satisfy the specifications of a parent class, a property known as *behavioral subtyping*[?]. Strong behavioral subtyping is a design principle in JML: any visible specification of a parent class is inherited by a derived class. Thus derived types inherit invariants from

Be sure we want this superseding design rather than merging – could use specs in the Java file if there is no JML declaration, just not merge when both have JML declarations.

¹Previous definitions of JML did require merging of specifications from multiple files; this requirement added complexity without appreciable benefit. The current design is simpler for tools, with the one drawback that the JML contents of a `.java` file is silently ignored when a `.jml` file is present, even if that `.jml` file does not contain a declaration of a particular entity.

their parent types and methods inherit behaviors from supertype methods they override.

For example, suppose method `m` in derived class `C` overrides method `m` in parent class `P`. In a context where we call method `m` on an object `o` with static type `P`, we will expect the specifications for `P.m` to be obeyed. However, `o` may have dynamic type `C`. Thus `C.m`, the method actually executed by the call `o.m()`, must obey all the specifications of `P.m`. `C.m` may have additional specifications, that is, additional behaviors, constraining its behavior further, but it may not relax any of the specifications given for `P.m`.

Specifications that are not visible in derived classes, such as those marked `private`, are not inherited, because a client cannot be expected to obey specifications that it cannot see. One additional exception to specification inheritance is method behaviors that are marked with the code modifier `??`. these behaviors apply only to the method of the class in which the behavior textually appears.

3.5 JML modifiers and Java annotations

The Java Modeling Language was defined prior to the introduction of annotations in Java. Some, but not all, of the features of JML can now be textually represented as Java annotations. Currently JML supports both the old and new syntactic forms.

3.5.1 Modifiers

Modifiers are JML keywords that specify JML characteristics of methods, classes, fields, or variables. Examples are `pure`, `model`, and `ghost`. They are syntactically placed just like Java modifiers, such as `public`.

Each such modifier has an equivalent Java annotation. For example

```
/*@ pure */ public int m(int i) ...
```

can be written equivalently as

```
org.jmlspecs.annotation.Pure public int m(int i) ...
```

The `org.jmlspecs.annotation` prefix can be made implicit in the usual way by including the import statement

```
import org.jmlspecs.annotation.Pure;
```

Note that in the second form, the `pure` designation is now part of the *Java* program and so the import of the `org.jmlspecs.annotation` package must also be in the Java program, and the package must be available to the Java compiler.

All of the modifiers, their corresponding Java annotations, and the locations in which they may be used are described in §20.3.5.

3.5.2 Type modifiers

Some modifiers are actually type modifiers. In particular `non_null` and `nullable` are in this category. Thus the description of the previous subsection (§20.1) apply to these as well.

However, Java 1.8 allows Java annotations to be applied to types wherever type names may appear. For example

```
(@NonNull String)toUpper(s)
```

is allowed in Java 1.8 but is forbidden in Java 1.7.

[Need additional discussion of the change in JML for Java 1.8, especially for arrays.](#)

3.6 Model and Ghost

[To be written](#)

3.7 Visibility

[To be written - note material written in Method Specifications section](#)

3.8 Evaluation and well-formedness of JML expressions

[To be written - note material written in Expressions chapter](#)

3.9 Null and non-null references

[To be written](#)

3.10 Static and Instance

[To be written](#)

3.11 Observable purity

To be written - perhaps this is a separate chapter

3.12 Location sets and Dynamic Frames

To be written - see section in DRM on Data Groups

3.13 Arithmetic modes

To be written - see later chapter - where shall we put this discussion

3.14 Immutable types and functions

To be written

3.15 Race condition detection

To be written - see later chapter - where shall we put this discussion

3.16 Redundant specifications

To be written

3.17 Controlling warnings

To be written - nowarn specifications - perhaps in the Syntax chapter; comment on possibility of unsoundness

3.18 `org.jmlspecs.lang` package

Some JML features are defined in the `org.jmlspecs.lang` package. The `org.jmlspecs.lang` package is included as a model import by default, just as the `java.lang` package is

imported by default in a Java file. `org.jmlspecs.lang.*` contains (at least²) these elements:

- `JML.informal(<string>)` : This method is a replacement for (and is equivalent to) the informal expression syntax `($??) (* ... *)`. Both expressions return a boolean value, which is always `true`.
- TBD

3.19 Interaction with other tools

3.19.1 Interaction with Type Annotations in Java 1.8

To be written

3.19.2 Interaction with the Checker framework

To be written

3.19.3 Interaction with FindBugs

To be written

²Tools implementing JML may add additional methods.

The conditional JML text feature presumes that the JML tool processing the JML text has a means to define various identifiers for the purpose of selecting or deselecting JML text. An identifier occurring before the @ in a given JML comment is *positive* if it is directly preceded by a + sign and negative if directly preceded by a - sign. The JML text in a conditional comment is included or ignored according to this rule.

- A conditional JML comment is included if (and only if) none of its negative identifiers (if any) are defined and there is at least one positive identifier that is defined.

Here are some examples, presuming the JML tool has defined the identifiers OPTA and OPTB, but not the identifier OPTC.

```
/*@ ... */ — included unconditionally
/**+OPTA@ ... */ — included because OPTA is defined and is positive
/**+OPTC@ ... */ — ignored because OPTC is not defined
/**-OPTA@ ... */ — ignored because OPTA is defined and but is negative
/**-OPTC@ ... */ — ignored because OPTC is not defined
/**+OPTA+OPTC@ ... */ — ignored because OPTA is defined and positive, even though OPTC is not defined
/**+OPTA-OPTC@ ... */ — ignored because OPTA is defined and positive, because though OPTC is negative it is not de
/**+OPTA-OPTB@ ... */ — ignored because OPTB is defined and is negative, even though OPTA is defined and positive
```

JML reserves the following identifiers:

- DEBUG : not defined by default; when defined, the JML debug statement is enabled
- ESC : tools should define this identifier when doing static checking
- RAC : tools should define this identifier when doing runtime checking

Issues with the JML textual format There are two issues that can arise with the syntactical design of JML. First, other tools may also use the @ symbol to designate comments that are special to that tool. If JML tools are trying to process files with such comments, the tools will interpret the comments as JML comments, likely causing a myriad of parsing errors.

Second, Java uses the @ sign to designate Java annotations. That in itself is not an ambiguity, but sometimes users will comment out such annotations with a simple preceding //, as in

```
//@MyAnnotation
```

This construction now looks like JML. The solution is to be sure there is whitespace between the // and the @, but it may not always be possible for the user to perform such edits.

4.2 JML Syntax

4.2.1 Syntax of JML specifications

JML specifications may be written as Java annotations. Currently these are only implemented for modifiers (cf. section TBD). In Java 8, the use of Java annotations for JML features will be expanded.

JML specifications may also be written in specially formatted Java comments: a JML specification includes everything between either (a) an opening `/*@` and closing `*/` or (b) an opening `//@` and the next line ending character (`\n` or `\r`) that is not within a string or character literal.

Such comments that occur within the body of a class or interface definition are considered to be a specification of the class, a field, or a method, depending on the kind of specification clause it is. JML specifications may also occur in the body of a method.

Obsolete syntax. In previous versions of JML, JML specifications could be placed within javadoc comments. Such specifications are no longer standard JML and are not supported by OpenJML.

4.2.2 Conditional JML specifications

JML has a mechanism for conditional specifications, based on a system of keys. A key is an identifier (consisting of ASCII alphanumeric and underscore characters, and beginning with a non-digit). A conditional JML comment is guarded by one or more positive or negative keys (or both). The keys are placed just before the `@` character that is part of the opening sequence of the JML comment (the `//@` or the `/*@`). Each key is preceded by a `'+'` or a `'-'` sign, to indicate whether it is a positive or negative key, respectively. *No white-space is allowed.* If there is white-space anywhere between the initial `//` or `/*` and the first `@` character, the comment will appear to be a normal Java comment and will be silently ignored.

The keys are interpreted as follows. Each tool that processes the Java+JML input will have a means (e.g. by command-line options) to specify the set of keys that are enabled.

- If the JML annotation has no keys, the annotation is always processed.
- If there are only positive keys, the annotation is processed only if at least one of the keys is enabled.
- If there are only negative keys, the annotation is processed unless one of the keys is enabled.
- If there are both positive and negative keys, the annotation is processed only if (a) at least one of the positive keys is enabled AND (b) none of the negative keys are enabled.

JML previously defined one conditional annotation: those that began with `/*+@` or `//+@`. ESC/Java2 also defined `/*-@` and `//-@`. Both of these are now deprecated. OpenJML does have an option to enable the `+-style` comments.

The particular keys do not have any defined meaning in the JML reference manual. OpenJML implicitly enables the following keys:

- **ESC** : the ESC key is enabled when OpenJML is performing ESC static checking;
- **RAC** : the RAC key is enabled when OpenJML is performing Runtime-Assertion-Checking.
- **DEBUG** : The DEBUG key is not implicitly enabled. However it is defined as the key that enables the **debug** JML statement. That is the **debug** statement is ignored by default and is used by OpenJML if the user enables the DEBUG key.
- **OPENJML** : The OPENJML key is enabled whenever OpenJML is processing annotations (and presumably is not enabled by other tools).
- **KEY**: The KEY key is reserved for annotations recognized by the KeY tool [?]. It is ignored by OpenJML.

Thus, for example, one can turn off a non-executable assert statement for RAC-processing but retain it for ESC and for type-checking by writing `//-RAC@ assert ...`

Chapter 5

JML Types

JML mostly uses Java types, but it does add a few additional types that may be used within JML specifications.

5.1 `\bigint`

The `\bigint` type is the set of mathematical integers (i.e., \mathbb{Z}). Just like Java primitive integral types are implicitly converted to `int` or `long`, all Java primitive integral types implicitly convert to `\bigint` where needed. When `\bigint` values need to be auto-boxed into an `Object`, they are boxed as `java.math.BigInteger` values; similarly when JML specifications are compiled for runtime checking, `\bigint` values are represented as `java.math.BigInteger` values. Within JML specifications, however, the `\bigint` type is treated as a primitive type.

5.2 `\real`

The `\real` type is the set of mathematical real numbers (i.e., \mathbb{R}). Just like the Java primitive types `float` is implicitly converted to `double`, both `real` and `double` values implicitly convert to `\real` where needed. When `\real` values need to be auto-boxed into an `Object`, they are boxed as `???TODO` values; similarly when JML specifications are compiled for runtime checking, `\real` values are represented as `???TODO` values. Within JML specifications, however, the `\real` type is treated as a primitive type.

5.3 \TYPE

TODO

Location set type? Object set type?

OLD STUFF:

5.3.1 JML types

Specifications are sometimes best written using infinite-precision mathematical types, rather than the fixed bit-width types of Java. JML's arithmetic modes (§??) allow choosing among various numerical precisions. In this section we simply note the type names that JML defines.

All of the Java type names are legal and useful in JML: `int` `short` `long` `byte` `char` `boolean` `double` `real` and class and interface types. In addition, JML defines the following:

- `\bigint` - the type of infinite-precision integers, represented as `java.lang.BigInteger` during run-time checking
- `\real` - the type of mathematical real numbers, represented as `TBD` during runtime-checking
- `\TYPE` - the type of JML type objects

The familiar operators are defined on values of the `\bigint` and `\real` types: unary and binary `+` and `-`, `*`, `/`, `%`. Also, these types can be used in quantified expressions and variables of these types can be declared as ghost or model variables.

The set of `\TYPE` values includes non-generic types such as `\type(org.lang.Object)`, fully parameterized generic types, such as `\type(org.utils.List<Integer>)`, and primitive types, such as `\type(int)`. The subtype operator (`<:`) is defined on values of type `\TYPE`.

TBD - what about other constructors or accessors of `\TYPE` values

Chapter 6

JML Specifications at Package and File level

There are no JML specifications at the package level. If there were, they would likely be written in `package-info.java` file. The only specifications JML specifications that are defined at the file level, applying to all classes defined in the file, are model import statements.

6.1 Model import statements

Java's import statements allow class and (with static import statements) field names to be used within a file without having to fully qualify them. The same import statements apply to names in JML annotations. In addition, JML allows *model import* statements. The effect of a JML model import statement is the same as a Java import statement, except that the names imported by the JML statement are only visible within JML annotations. If the model import statement is within a .jml file, the imported names are visible only within annotations in the .jml file, and not outside JML annotations and not in a corresponding .java file. These are import statements that only affect name resolution within JML annotations and are ignored by Java. They have the form

```
//@ model <Java import statement>;
```

Note that both

```
model <Java import statement>;
```

and

```
/*@ model */<Java import statement>;
```

are invalid. The first is not within a JML comment and is illegal Java code. The second is a normal Java import with a comment in front of it that would have no effect in JML, even if JML recognized it (tools should warn about this erroneous use).

6.2 Default imports

The Java language stipulates that `java.lang.*` is automatically imported into every Java compilation unit. Similarly in JML there is an automatic model import of `org.jmlspecs.lang.*`. However, there are not yet any standard-defined contents of the `org.jmlspecs.lang` package.

is it called a compilation unit?
Is this correct?

6.3 Issues with model import statements

As of this writing, no tools distinguish between Java import statements and JML import statements.

SAY MORE

Chapter 7

Specifications of Java types

Need to work out specs for enum types

Specifications at the class level serve three different primary purposes: specifications that are applied to all methods in the class, specifications that state properties of the data structures in the class, and declarations that help with information hiding.

7.1 `non_null_by_default`, `nullable_by_default`, `@NonNullByDefault`, `@NullableByDefault`

The `non_null_by_default` and `nullable_by_default` modifiers or, equivalently, the `@NonNullByDefault` and `@NullableByDefault` Java annotations, specify the default nullity declaration within the class. The default applies to all field and local variable declarations, to formal parameters and method return values, and recursively to any nested or inner classes that do not have nullity declarations of their own.

7.2 invariant clause

TODO

7.3 constraint clause

TODO

7.4 initially clause

TODO

7.5 ghost fields

TODO

7.6 model fields

TODO

7.7 represents clause

TODO

7.8 model methods and model classes

TODO

7.9 initializer and static_initializer

TODO

7.10 axiom

TODO

7.11 readable if clause and writiable if clause

TODO

7.12 monitors_for clause

TODO

Chapter 8

JML Method specifications

Method specifications describe the behavior of the method. JML is a modular specification methodology, with the Java method being the fundamental unit of modularity. The specifications may under-specify a method. For example, the specifications may simply say that the method always returns normally (with throwing an exception), but give no constraints on the value returned by the method. The degree of precision needed will depend on the context.

8.1 Structure of JML method specifications

```
<method-spec> ::= ( also )? <behavior-seq>
                  ( also implies_that <behavior-seq> )?
                  ( also for_example <behavior-seq> )?

<behavior-seq> ::= <behavior> ( also <behavior> )*

<behavior> ::=
    ( ( <java-visibility> ( code )? <behavior-id> )? <clause-seq> )
    | <java-visibility> ( code )? <model-program>

<java-visibility> ::= ( public | protected | private )?

<behavior-id> ::= behavior | normal_behavior | exceptional_behavior
                 | behaviour | normal_behaviour | exceptional_behaviour

<clause-seq> ::= ( <clause> | <nested-clause> )+

<clause> ::= <requires-clause>
             | <assignable-clause>
             | <ensures-clause>
             | <signals-clause>
             | <signals-only-clause>
```



```

| <diverges-clause>
| <measured-by-clause>
| <duration-clause>
| <when-clause>
| <old-clause>
| <forall-clause>
| <working-space-clause>
| <accessible-clause>
| <callable-clause>
| <captures-clause>

<nested-clause> ::= { | <clause-seq> ( also <clause-seq> ) * | }

```

Meta-parser rules:

- Each of the behavior keywords spelled behaviour is equivalent to the corresponding keyword spelled behavior. The latter is more common.
- A behavior beginning with normal_behavior may not contain a <signals-clause> or a <signals-only-clause>. It implicitly contains the clause **signals (Exception e) false;**.
- A behavior beginning with exceptional_behavior may not contain a <ensures-clause>. It implicitly contains the clause **ensures false;**.
- Clause ordering: in any consecutive sequence of «>clause», any <old-clause> and <forall-clause> must appear before any <requires-clause>, which must appear before any other clauses. *Is this a requirement or a style recommendation?*
- *May any clauses appear after a nested-clause?*

Note that the vertical bars in the production for *nested-clause* are literals, not meta-symbols.

OK to relax the rules on which clauses can be present where in parsing, and then enforce or advise in later checking?

FIXME - the method-spec production is not correct - the first also might be omitted whichever one it is

8.1.1 Behaviors

The basic structure of JML method specifications is as a set of *behaviors*. Each behavior contains a set of *clauses*. The various kinds of clauses are described in the subsequent sections of this chapter. Each kind of clause has a default, which applies if the clause is textually absent from the behavior. A particular kind of clause is the

requires clause, which gives the *precondition* for the behavior; the other clauses state the properties that are true about the method's execution.

For each behavior, if the method is called in a context in which the behavior's precondition is true, then the method must adhere to the constraints specified by the remaining clauses of the behavior. Not all of the preconditions need be true, but at least one of them must be, otherwise the method is being called in a context in which its behavior is undefined. For example, a method's specification may have two behaviors, one with a precondition that the method's argument is not null, and the other behavior with a precondition that the method's argument is null. In this case, in any context, one or the other behavior will be active. If however, the second behavior were not specified, then it would be a violation to call the method in any context other than those in which the first precondition, that the argument is not null, is true. More than one behavior may be active (have its precondition true); every active behavior must be obeyed by the method. Where preconditions are not mutually exclusive, care must be taken that the behaviors themselves are not contradictory, or it will not be possible for any implementation to satisfy the combination of behaviors.

8.1.2 Nested specification clauses

Nested specification clauses are syntactic shorthand for an expanded equivalent in which clauses are replicated. The nesting syntax simply allows expressing common subsequences of clauses to be expressed without repetition, where that helps understandability.

In particular, referring to the grammar above, a *<behavior>* whose *<clause-seq>* contains a *<nested-clause>* is equivalent to a sequence of *<behavior>*s as follows:

if *<nested-clause>_A* is a combination of *n* *<clause-seq>* as in

{ | *<clause-seq>_{S1}* (**also** *<clause-seq>_{Si}*) * | }

then

(*<java-visibility>_V* (**code**) ? *<behavior-id>_X*) ? *<clause>*_D* *<nested-clause>_A* *<clause-seq>_E*

is equivalent to a sequence of *n* *<behavior>* constructions

(*<java-visibility>_V* (**code**) ? *<behavior-id>_X*) ? *<clause>*_D* *<clause-seq>_{S1}* *<clause-seq>_E*

also

...

also

(*<java-visibility>_V* (**code**) ? *<behavior-id>_X*) ? *<clause>*_D* *<clause-seq>_{Sn}* *<clause-seq>_E*

Is there a better way to describe this desugaring? and a better way to format it?

8.1.3 Specification Inheritance

§3.4 Briefly restate specification inheritance, code modifier, details of when the spec applies

8.1.4 Modifiers and annotations

TODO

8.1.5 Visibility

The following discussion has some errors and needs fixing; also need to talk about `spec_public`, `spec_protected`

Each method specification behavior has a *java-visibility* (cf. the discussion in §??). Any of the kinds of behavior keywords (`behavior`, `normal_behavior`, `exceptional_behavior`) may be prefixed by a Java visibility keyword (`public`, `protected`, `private`); the absence of a visibility keyword indicate package-level visibility. A lightweight behavior (one without a behavior keyword) has the visibility of its associated method.

The visibility of a behavior determines the names that may be referenced in the behavior. The general principle is that a client that has permission to see the behavior must have permission to see the entities in the behavior. Thus

any name (of a type, method or field) in a method specification visible to a client must also be visible to the client.

For example, a public behavior may contain only public names. A private behavior may contain any name visible to a client that can see the private names; this would include other private entities in the same or enclosing classes, any public name, any protected name from super classes, and any package or protected name from other classes in the same package. The visibility for protected and package behaviors is more complex. A protected behavior is visible to any client in the same or subclasses; since the subclasses may be in a different package, the protected behavior may contain other names with protected visibility only if they are visible in the behavior by virtue of inheritance, and not if they are visible only because of being in the same package. To be explicit, suppose we have class A, unrelated class B in the same package, class C a superclass of A in a different package, and class D derived from A but in a different package, with identifiers A.a, B.b, and C.c each with protected visibility. Only A.a and C.c are visible in class D; thus a protected behavior in class A, which is visible to D, may contain A.a and C.c but not B.b. Similarly a behavior with package visibility may only contain names that are visible by virtue of being in the same package (and public names); names with protected visibility that are visible in a class by virtue of inheritance are not necessarily visible to clients who can see the package-visible behavior.

The root of the complexity is that protected visibility is not transitive, whereas the other kinds of Java visibility are. Conceptually, protected visibility must be separated into two kinds of visibility: protected-by-inheritance and protected-by-package. Each of these is separately transitive. Then the visibility rules can be summarized in Table 8.1.

Table 8.1: Visibility rules for method specification behaviors

Behaviors with this visibility	may contain names that are visible in the class because of this visibility
public	public
protected	public, protected-by-inheritance
package	public, protected-by-package, package
private	any

8.1.6 Grammar of method specifications

Fillin – remember lightweight, behavior, normal_behavior, exceptional_behavior, examples, implies_that, visibility, model program behaviors, also, nested behaviors

Do we relax the ordering and the constraints on nesting that are in the current Ref-Man

Comment on comparison with ACSL

8.2 Method specifications as Annotations

8.3 Modifiers for methods

TODO

8.4 Common JML method specification clauses

TODO

8.4.1 requires clause

TODO

8.4.2 ensures clause

TODO

8.4.3 assignable clause

TODO

8.4.4 signals clause

TODO

8.4.5 signals_only clause

TODO

8.5 Advanced JML method specification clauses

TODO

8.5.1 accessible clause

TODO

8.5.2 diverges clause

TODO

8.5.3 measured_by clause

TODO

8.5.4 when clause

TODO

8.5.5 old clause

TODO

8.5.6 forall clause

TODO

8.5.7 duration clause

TODO

8.5.8 working_space clause

TODO

8.5.9 callable clause

TODO

8.5.10 captures clause

TODO

8.6 Model Programs (model_program clause)**8.6.1 Structure and purpose of model programs****8.6.2 extract clause**

TODO

8.6.3 choose clause

TODO

8.6.4 choose_if clause

TODO

8.6.5 `or` clause

TODO

8.6.6 `returns` clause

TODO

8.6.7 `continues` clause

TODO

8.6.8 `breaks` clause

TODO

8.7 TODO Somewhere

constructor field method nowarn

<:: token

lots more backslash tokens

Chapter 9

Field Specifications

Fields may have various modifiers, each of which states a restriction on how the field may be used. Fields may be part of *data groups*, which allow specifying frame conditions on fields that may not be visible because of the Java visibility rules. Also, a specification may introduce *ghost* or *model* fields that are used in the specification but are not present in the Java program.

9.1 Field and Variable Modifiers

The modifiers permitted on a field, variable, or formal parameter declaration are shown in Table 9.1.

9.1.1 non_null and nullable (@NonNull, @Nullable)

The non_null and nullable modifiers, and equivalent @NonNull and @Nullable annotations, specify whether or not a field, variable, or parameter may hold a null value. The modifiers are valid only when the type of the modified construct is either a reference or array type, not a primitive type.

Discuss @Non-Null TestJava a, b;
Need to discuss relationship with JSR308

9.1.2 spec_public and spec_protected (@SpecPublic, @SpecProtected)

These modifiers are used to change the visibility of a Java field when viewed from a JML construct. A construct labeled spec_public has public visibility in a JML specification, even if the Java visibility is less than public; similarly, a construct labeled spec_protected has protected visibility in a JML specification, even if the Java

Table 9.1: Modifiers allowed on field, variable and parameter declarations

Modifier	Where	Purpose
<code>non_null</code>	field, var, param	the variable may not be null (§9.1.1)
<code>nullable</code>	field, var, param	the variable may be null
<code>spec_public</code>	field	visibility is public in specs
<code>spec_protected</code>	field	visibility is protected in specs
<code>model</code>	field	representation field
<code>ghost</code>	field, var	specification only field
<code>uninitialized</code>	var	TBD
<code>instance</code>	field	not static
<code>monitored</code>	field	guarded by a lock
<code>secret</code>	field, var, param	hidden field
<code>peer</code>	field, param	TBD
<code>rep</code>	field, param	TBD
<code>readonly</code>	field, param	TBD

visibility is less than protected. Section ?? contains a detailed discussion of the effect of information hiding using Java visibility on JML specifications.

Listing 9.1: Use of `spec_public`

```

private /*@ spec_public */ int value;

/*@ ensures value == i;
public setValue(int i) {
    value = i;
}

```

For example, Listing 9.1 shows a simple setter method that assigns its argument to a private field named `value`. The visibility rules require that the specifications of a public method (`setValue`) may reference only public entities. In particular, it may not mention `value`, since `value` is private. The solution is to declare, in JML, that `value` is `spec_public`, as show in the Listing.

9.1.3 `ghost` **and** `@Ghost`

9.1.4 `model` **and** `@Model`

9.1.5 `uninitialized` **and** `@Uninitialized`

9.1.6 `instance` **and** `@Instance`

9.1.7 `monitored` **and** `@Monitored`

9.1.8 `query, secret` **and** `@Query, @Secret`

9.1.9 `peer, rep, readonly` (`@Peer, @Rep, @Readonly`)

Check `readonly`
vs. `read_only`,
`Readonly` vs.
`ReadOnly`

9.2 Datagroups: `in` and `maps` clauses

[TODO](#)

9.3 Ghost fields

[TODO](#)

9.4 Model fields

[TODO](#)

Chapter 10

JML Statement Specifications

JML Statement specifications are JML constructs that appear as statements within the body of a Java method or initializer. Some are standalone statements, while others are specifications for loops or blocks that follow.

Grammar:

```
<jml-statement> ::=
    <jml-assert-statement>
    | <jml-assume-statement>
    | <jml-local-declaration>
    | <jml-unreachable-statement>
    | <jml-reachable-statement>
    | <jml-set-statement>
    | <jml-debug-statement>
    | <jml-loop-specification>
    | <jml-hence-by-statement>
    | <jml-refining-specification>
```

10.1 assert statement and Java assert statement

Grammar:

```
<jml-assert-statement> ::= assert <expression> ;
```

Type checking requirements:

- the <expression> must be boolean

The `assert` statement that the given expression must be true at that point in the program. A static checking tool is expected to require a proof that the asserted expression is true and to issue a warning if the expression is not provable. A runtime assertion checking tool is expected to check whether the asserted expression is true and to issue a warning message if it is not true in the given execution of the program.

Note that a JML `assert` statement has a different effect than a Java `assert` statement. When assertion checking is enabled, a Java `assert` statement will result in a `AssertionError` at runtime if the corresponding assertion evaluates to false; if assertion checking is disabled (the default), a Java `assert` statement is ignored. Runtime assertion checking tools may implement JML `assert` statements as Java `assert` statements.

10.2 `assume` statement

Grammar:

`<jml-assume-statement> ::= assume <expression> ;`

Type checking requirements:

- the `<expression>` must be boolean

The `assume` statement adds an assumption that the given expression is true at that point in the program.

Static analysis tools may assume the given expression to be true. Runtime assertion checking tools may choose to check or not to check the `assume` statements.

An `assume` statement might be used to state an axiom or fact that is not easily proved. However, `assume` statements should be used with caution. Because they are assumed but not proven, if they are not actually true an unsoundness will be introduced into the program. For example, the statement `assume false;` will render the following code silently infeasible. Even this may be useful, since, during debugging, it may be helpful to shut off consideration of certain branches of the program.

10.3 `ghost` and `model` declarations

10.4 `unreachable` statement

Grammar:

`<jml-unreachable-statement> ::= unreachable <expression>? ;`

Type checking requirements:

- the optional `<expression>` must be boolean; if not present its default value is **true**.

The `unreachable` statement asserts that no feasible execution path will ever reach this statement when the given expression is true.

The terminating semicolon is required and is easily forgotten, since statement is almost always used without the optional expression.

The `unreachable` statement with an expression is an extension to standard JML.

It has been common practice to insert `assert false;` statements to check whether a given program point is infeasible. The `unreachable` statement accomplished the same purpose with clearer syntax.

10.5 reachable statement

Grammar:

`<jml-reachable-statement> ::= reachable <expression>? ;`

Type checking requirements:

- the optional `<expression>` must be boolean; if not present its default value is **true**.

The `reachable` statement asserts that there exists a feasible execution path that reaches this statement with the given expression true.

The terminating semicolon is required and is easily forgotten, since statement is almost always used without the optional expression.

The usual execution of the underlying solver checks whether there are any feasible paths for which an assertions are false. The reachability test is different and typically requires separate executions of underlying solvers, as the test is now for at least one path that reaches the given statement.

The `reachable` statement is an extension to standard JML.

It has been common practice to insert `assert false;` statements to check whether a given program point is feasible; if it is feasible, the `assert false;` statement will cause a static checking warning. The `reachable` statement accomplished the same purpose with clearer syntax.

10.6 set and debug statements

Grammar:

```
<jml-set-statement> ::= set <statement>  
<jml-debug-statement> ::= debug <statement>
```

Type checking requirements:

- the <statement> may be any single statement, including a block statement

The DRM requires a set statement to take an assignment expression; the DRM is inconsistent in how it describes debug statements.

10.7 statement (block) specification

[needs index entry](#)

[Write this](#)

10.8 loop specifications

[needs index entry](#)

[Write this](#)

Chapter 11

JML Expressions

JML Expressions can include most of the operations defined in Java and additional operations defined only in JML. JML operations are one of four types:

- infix operations that use non-alphanumeric symbols (e.g., `<==>`)
- identifiers that begin with a backslash (e.g., `\result`)
- identifiers that begin with a backslash but have a functional form (e.g., `\old`)
- methods defined in JML whose syntax is Java-like (e.g., `JML.informal(...)`)

The Java-like forms replicate some of the backslash forms. The backslash forms are traditional JML and more concise. However, the preference for new JML syntax is to use the Java-like form since supporting such syntax requires less modification of JML tools.

11.1 `org.jmlspecs.lang.JML`

[Say more](#)

11.2 Java operations used in JML

Specification expressions must not have side effects. During run-time assertion checking, the execution of specifications may not change the state of the program under test. Thus some Java operators are not permitted in JML expressions:

- allowed: `+ - * / % == != <= >= < > . ^ & | && || << >> >>> ? :`
- prohibited: `++ -- = += -= *= /= %= &= |= ^= <<= >>= >>>=`

Table 11.1: Java and JML precedence (cf. <https://docs.oracle.com/javase/tutorial/java/nutsandbolts/operators.html>)

Java operator	JML operator	
highest precedence		associativity
literals, names, parenthesis	quantified	left
postfix: <code>.</code> <code>[]</code> method calls		right
prefix: unary <code>+</code> <code>-</code> <code>!</code> <code>~</code> cast new		left
<code>*</code> <code>/</code> <code>%</code>		left
binary <code>+</code> <code>-</code>		left
<code><<</code> <code>>></code> <code>>>></code>		left
<code><=</code> <code><</code> <code>>=</code> <code>></code> <code>instanceof</code>	<code><:</code> <code><#</code> <code><#=</code>	left
<code>==</code> <code>!=</code>		left
<code>&</code>		left
<code>^</code>		left
<code> </code>		left
<code>&&</code>		left
<code> </code>		left
	<code>==></code> <code><==</code>	right
	<code><==></code> <code><!=></code>	left
<code>?:</code>		right
assignment, assign-op		right
lowest precedence		

11.3 Precedence of infix operations

JML infix operators may be mixed with Java operators. The new JML operators have precedences that fit within the usual Java operator precedence order, as shown in Table 11.1.

fix complement
operator

11.4 Well-defined expressions

An expression used in a JML construct must be well-defined, in addition to being syntactically and type-correct. This requirement disallows the use of partial functions with arguments for which the result of the function is undefined. For example, the expression $(x/0) == (x/0)$ is considered in JML to be not well-defined (that is, undefined), rather than true by identity. An expression like $(x/y) == (x/y)$ (for integer x and y) is true if it can be proved that y is not 0. For example, $y \neq 0 \implies ((x/y) == (x/y))$ is well-defined and true.

The well-definedness rules for JML operators are given in the section describing that operator. The rules for Java operators are given here. They presume that the expressions are type correct.

- literals and names are well-defined
- parenthesis operator: well-defined iff the operand is well-defined
- member-of (dot operator): well-defined iff the expression to the left of the dot is well-defined and not null
- array element (`[]` operator): well-defined iff the array expression is well-defined and not null and the index expression is well-defined and within range
- method calls: well-defined iff (a) the receiver and all arguments are well-defined and (b) if the method is not static, the receiver is not null and (c) the method's precondition is true and (d) the method does not throw any Exceptions
- cast operator: well-defined iff the argument is well-defined and, for limited range integral types, the result is not out of range for the type
- new operator: well-defined iff all arguments to the constructor call are well-defined
- bit operators and non-short-circuit boolean operators (`&` `|` `^` `~`): well-defined iff all operands are well-defined
- shift operators (`<<` `>>` `>>>`): well-defined iff all operands are well-defined. Note that Java defines the shift operations for any value of the right-hand operand; the value is trimmed to 5 or 6 bits by a modulo operation appropriate to the bit-width of the left-hand operand. JML tools may choose to raise a warning if the value of the right-hand operand is outside the 'expected' range.
- divide and modulo operators: well-defined iff both operands are well-defined and the divisor is not zero and, for limited range integral types, the result is not out of range for the type
- other arithmetic operators: well-defined iff both operands are well-defined and, for limited range integral types, the result is not out of range for the type
- string concatenation: well-defined if the operands are well-defined (they may be null)
- boolean negation operator (`!`): well-defined iff the operand is well-defined
- relational and equality operators: well-defined iff both operands are well-defined
- short-circuit boolean and operator (`&&`): well-defined iff
 - the left operand is well-defined and
 - either the left operand is false or the right-operand is well-defined
- short-circuit boolean or operator (`||`): well-defined iff
 - the left operand is well-defined and
 - either the left operand is true or the right-operand is well-defined
- ternary operator (`?:`): well-defined iff
 - the condition is well-defined and
 - the then operand is well-defined whenever the condition is true and
 - the else operand is well-defined whenever the condition is false.

floating point operations?

For example `(o != null ? o.x : 0)` is well-defined (if it is type-correct) because (a) the condition `o != null` is well-defined and (b) the then expression `o.x` is well-defined if the condition `o != null` is true and (c) the else expression is always well-defined.

11.5 Implies and reverse implies: \implies \impliedby

Type information:

- two arguments, each an expression of boolean type
- a \implies expression is well-defined if the left operand is well-defined and either the left operand is false or the right operand is well-defined
- a \impliedby expression is well-defined if the left operand is well-defined and either the left operand is true or the right operand is well-defined
- result is boolean

The \implies operator denotes implication and is a short-circuit operator. It is true if the left-hand operand is false or the right-hand operand is true; if the left operand is false, the right operand is not evaluated and may be undefined. The operation

$$\langle left \rangle \implies \langle right \rangle$$

is equivalent to

$$(! \langle left \rangle) \vee \langle right \rangle$$

The \impliedby operator denotes reverse implication and is a short-circuit operator. It is true if the left-hand operand is true or the right-hand operand is false; if the left operand is true, the right operand is not evaluated and may be undefined. The operation

$$\langle left \rangle \impliedby \langle right \rangle$$

is equivalent to

$$\langle left \rangle \vee (! \langle right \rangle)$$

Both operators are right associative. For example $P \implies Q \implies R$ is parenthesized as $P \implies (Q \implies R)$. This is the natural association from logic: $(P \implies Q) \implies R$ is equivalent to $(P \ \&\& \ !Q) \vee R$ whereas $P \implies (Q \implies R)$ is equivalent to $!P \vee !Q \vee R$.

Note that because of the short-circuit nature of these two operators $\langle left \rangle \implies \langle right \rangle$ is not necessarily equivalent to $\langle right \rangle \impliedby \langle left \rangle$. In particular, if $\langle right \rangle$ is undefined then $\langle left \rangle \implies \langle right \rangle$ is either true or undefined, whereas $\langle right \rangle \impliedby \langle left \rangle$ is always undefined.

11.6 Equivalence and inequivalence: \iff $\impliedby \implies$

Type information:

- two arguments, each an expression of boolean type
- the expression is well-defined if both operands must be well-defined
- result is boolean

The \iff operator denotes equivalence: it is true if both operands are true or both are false. It is equivalent to equality (\equiv), except that it is lower precedence. For example, $P \ \&\& \ Q \iff R \vee S$ is $(P \ \&\& \ Q) \iff (R \vee S)$, whereas $P \ \&\& \ Q \equiv R \vee S$ is $(P \ \&\& \ (Q \equiv R)) \vee S$.

The $\lt!=\gt$ operator denotes inequivalence: it is true if one operand is true and the other false. It is equivalent to inequality (\neq), except that it is lower precedence. For example, $P \ \&\& \ Q \ \lt!=\gt \ R \ || \ S$ is $(P \ \&\& \ Q) \ \lt!=\gt \ (R \ || \ S)$, whereas $P \ \&\& \ Q \ != \ R \ || \ S$ is $(P \ \&\& \ (Q \ != \ R)) \ || \ S$.

Both of these operators are associative and commutative. Accordingly left- and right-associativity are equivalent. The operators are not chained: $P \ \lt==\gt \ Q \ \lt==\gt \ R$ is $(P \ \lt==\gt \ Q) \ \lt==\gt \ R$, not $(P \ \lt==\gt \ Q) \ \&\& \ (Q \ \lt==\gt \ R)$; for example, $P \ \lt==\gt \ Q \ \lt==\gt \ R$ is true if P is true and Q and R are false. Similarly $P \ \lt!=\gt \ Q \ \lt!=\gt \ R$ is $(P \ \lt!=\gt \ Q) \ \lt!=\gt \ R$ and is true if P is true and Q and R are false.

11.7 JML subtype: $\lt:$

Type information:

- two arguments, each of type $\backslash\text{TYPE}$
- both operands must be well-defined
- result is boolean

The $\lt:$ operator denotes JML subtyping: the result is true if the left operand is a subtype of the right operand. Note that the argument types are

TYPE , that is JML types (cf. §??). [Say more about relationship to Java subtyping](#)

Note that the operator would be better named $\lt:=$, since it is true if the two operands are the same type.

11.8 Lock ordering:

Type information:

- two arguments, each of reference type
- both operands must be well-defined and both must not be null
- result is boolean

[Say more](#)

11.9 $\backslash\text{result}$

Type information:

- no arguments
- always well-defined, if type-correct

- result type is the return type of the method in whose specification the expression appears
- may only be used in `ensures` clauses

The `\result` expression denotes the value returned by a method. The expression is only permitted in `ensures` clauses of the method's specification. It is a type-error to use `\result` in the specification of a constructor or a method whose return type is `void`.

11.10 `\exception`

Type information:

- no arguments
- always well-defined, if type-correct
- result type is the type of the exception given in the `signals` clause
- only permitted in the `signals` clause

`\exception` is an Open-JML extension

The `\exception` expression denotes the exception object within a `signals` clause. The expression is only permitted in `signals` clause of the method's specification. It is an alternative form to using a variable declared in the `signals` clause's declaration. For example, the following two constructions are equivalent:

```
//@ signals (RuntimeException e) ... e ... ;
//@ signals (RuntimeException) ... \exception ... ;
```

11.11 `\old`, `\pre`, and `\past`

`\past` is an extension

Text needed

11.12 `\lblpos`, `\lblneg`, `\lbl`, and `JML.lblpos()`, `JML.lblneg()`, `JML.lbl()`

`\lbl` and the JML forms are extensions

Type information: `\lblpos`, `\lblneg`

- special syntax, including a boolean expression
- well-defined if the boolean expression is well-defined
- result type is boolean; value is the same as the expression in the argument

Type information: `JML.lblpos()`, `JML.lblneg()`

- first argument is a String literal, the second has boolean type
- well-defined if the arguments are well-defined
- result type is boolean; value is the second argument

Type information: `\lbl`

- special syntax, including an expression
- well-defined if the argument expression is well-defined
- result type is the same as the expression; value is the value of the expression

Type information: `JML.lbl()`

- first argument is a String literal, the second has any type
- well-defined if the arguments are well-defined
- result type is the same as the type of the second argument; value is the second argument

These expressions are used for debugging. When static checking finds that some assertion is invalid and generates a counterexample for that invalid assertion, then the value of the expression in contained in the `lbl` expression is reported as part of the counterexample. When runtime assertion checking is used, these expressions print the String `id` and the value of expression. In each case, the construct simply passes on the type and value of its expression, possibly generating some debug output in the process.

The `\lblpos`, `\lblneg`, and `\lbl` expressions have a non-standard syntax:

(`\lbl` *<id>* *<expression>*)

with no comma between what would be the arguments. Here the *<id>* is a Java identifier. The `JML.lbl` forms are standard functional forms: the first argument is a String literal; the second is an expression. A String literal is required instead of a String expression because the value is used to identify the output as coming from this expression and it is not evaluated during static checking.

In the positive and negative forms, the argument is a boolean expression. Output is generated by `lblpos` only if the expression is true; output is generated by `lblneg` only if the expression is false. In the neutral form (`\lbl` and `JML.lbl`), output is generated whatever the value. For any type, the value is converted to a String value as is customary in Java (e.g., using `toString()`).

Examples: [Examples needed](#)

11.13 \nonnull elements

Type information:

- strict JML: one argument, an expression of array type, may be null
- OpenJML extensions: any number of arguments, each an expression of array type
- always well-defined, if type-correct
- result type is boolean

The argument of the `\nonnullElements` expression must be an expression that evaluates to an array. The `\nonnullElements` expression is true if the argument is true and each element of the array (if any) is not null.

If there are multiple arguments, then the result is true if all of the arguments are non-null and have all non-null array elements. If `\nonnullElements` has no arguments, its value is true.

Perhaps allow the argument to be reference type - result is true if the argument is non-null and, if an array, all elements are non-null. Are elements non-null recursively?

Allow any number of arguments?

11.14 `\fresh`

Type information:

- strict JML: one argument, an expression of reference type
- OpenJML extensions: any number of arguments, each an expression of reference type
- well-defined, if type-correct and all arguments are non-null
- result type is boolean
- use: `\fresh` may be used only in `ensures` and `signals` clauses

The argument of the `\fresh` expression must be an expression that evaluates to a non-null reference. The `\fresh` expression is true if the argument is a reference that was not allocated in the pre-state.

If there are multiple arguments, then the result is true if each argument is itself `\fresh`; that is, the value with n arguments is the conjunction of n terms each of which is `\fresh` applied to one of the arguments in turn. If `\fresh` has no arguments, its value is true.

Standardize the JML extension?

11.15 informal expression: (*...*) and JML.informal()

Type information:

- one argument
- always well-defined, if type-correct
- result type is boolean; value is always true

The syntax of the informal expression is

(* ... *) ,

where the ... denotes any sequence of characters not including line breaks or the two-character sequence *). An alternate form is

JML.informal(<expression>)

where <expression> is a String expression, though typically a String literal. The character sequence and the string expression are natural language text that may be ignored by JML tools; the intent is to convey to the reader some natural language specification that will not be checked by automated tools.

In the second form, the argument is type checked and must have type `java.lang.String`; it is not evaluated. It is generally a string literal.

The expression always has the value true.

Examples:

```
//@ ensures (* data structure is self-consistent *);
//@ ensures JML.informal("data structure is self-consistent");
public void m() ...
```

11.16 \type

Type information:

- one argument, a type name
- always well-defined, if type-correct
- result type is \TYPE

This expression is a type literal. The argument is the name of a type as might be used in a declaration; the type may be a primitive type, a non-generic reference type, a generic type with type arguments or an array type. The value of the expression is the JML type value corresponding to the given type. It is analogous to `.class` in Java, which converts a type name to a value of type `Class`. The type name is resolved like any other type name, with respect to whatever type names are in scope.

Generic types must be fully parameterized; no wild card designations are permitted.

What about type variables

For more discussion of JML types and their relationships to Java type, see §??.

Examples:

```
//@ ... \type(int) ...
//@ ... \type(Integer) ...
//@ ... \type(java.lang.Integer) ...
//@ ... \type(java.util.LinkedList<String>) ...
//@ ... \type(java.util.LinkedList<String>[]) ...
```

11.17 \typeof

Type information:

- one expression argument, of any type
- well-defined if the argument is well-defined and not null
- result type is \TYPE

The \typeof expression returns the dynamic type of the expression that is its argument. In run-time checking this may require evaluating the argument. This operation returns a JML type (\typeof); it is analogous to the Java method `.getClass()`, which returns a Java type value (of type `Class`).

Verify that primitive types are allowed

Examples:

```
Object o = new Integer(5);
// o has static type Object, but dynamic type Integer
//@ assert \typeof(o) == \type(Integer); // - true
//@ assert \typeof(o) == \type(Object); // - false
//@ assert \typeof(5) == \type(int); // - true
```

11.18 \elemtype

Type information:

- one argument, of type \TYPE
- well-defined iff the argument is well-defined and the value is an array type
- result type is \TYPE

This operator returns the element type of an array type.

What if the argument is not an array type? Should we allow a null value of `\TYPE`

Examples:

```
//@ assert \elemtype(\type(int[])) == \type(int);
//@ assert \elemtype(\type(int)) == \type(int); // - undefined
```

11.19 `\is_initialized`

Text needed

11.20 `\invariant_for`

Type information:

- Strict JML: one argument, of type `Object` or a typename, but not a primitive expression
- Extension: any number of arguments
- well-defined iff the argument is well-defined and the value is a reference type but not an array type
- result type is `boolean`

This expression with one expression argument is equivalent to the conjunction (with `&&`) of the static and non-static JML-visible invariants in the static type of the receiver and all its super classes and interfaces (recursively), with the argument as the receiver for the invariants.

OpenJML allows as an extension one typename argument. The expression with a typename argument is equivalent to the conjunction (with `&&`) of the *static* JML-visible invariants in the named type and all its super classes and interfaces (recursively).

In each case the order of invariants is (1) that invariants of super classes and interfaces occur before derived classes and interfaces, (2) `Object` is first and the named type is last, and (3) within a type, invariants occur in textual order.

Strict JML requires that there be just one argument. As an extension, OpenJML allows any number of arguments. The expression is then equivalent to the conjunction of the values for each argument, in order, conjoined by the short-circuit operator `&&`. When there are no arguments, the value of `\invariant_for()` is `true`.

Are the invariants of superclasses and interfaces included?

Are both static and non-static invariants included?

The DRAFT JML reference manual does not mention JML-visibility, but I presume that must be the case.

multiple arguments, type-name argument

The DJMLRM does not mention a static-only version of `\invariant_for` - so this is an extension?

11.21 `\not_modified`

Type information:

- Strict JML: one argument, an expression of any type
- Extension: any number of arguments, expressions of any type
- well-defined
- result type is boolean
- use: only in `ensures` or `signals` clauses

A `\not_modified` expression is a two-state expression that may occur only in `ensures` or `signals` clauses. It satisfies this equivalence:

$$\text{\not_modified}(o) == (\text{\old}(o) == (o))$$

The argument may be null.

A `\not_modified` expression with multiple arguments is the conjunction of the corresponding terms each with one argument; if `\not_modified` has no arguments, its value is true.

The RM says the argument is a store-ref list, rather than an expression. Which do we want? A store-ref-list allows constructions such as `o.*` or `a[*]` or `a[1..6]` but not `a+b`.

11.22 `\lockset` and `\max`

Text needed

11.23 `\reach`

Text needed

11.24 `\duration`

Text needed

11.25 `\space`

Text needed

11.26 `\working_space`

Text needed

not-assigned, not-modified, only-accessed, only-called, only-assigned, only-captured,
spec-quantified-expr Text needed

Chapter 12

Arithmetic modes

TODO

Chapter 13

Universe types

Chapter 14

Model Programs

Describe the intent, syntax and semantics of model programs

Chapter 15

Obsolete and Deprecated Syntax

describe deprecated syntax, obsolete syntax, incompatible changes from past versions

Chapter 16

Grammar Summary

Automatic collection of all of the grammar productions listed elsewhere in the document

Chapter 17

Type Checking Summary

This was in the DRM outline - is there something to be put in here? If it is to be collected from the rest of the document, we need to place markers to identify the relevant stuff.

Chapter 18

Verification Logic Summary

This was in the DRM outline. What was its intent? Is it the same as a section on semantics and translation?

Chapter 19

Differences in JML among tools

Some material is in the DRM. Needs to be enhanced. Should have a detailed comparison with ACSL, for example – see the appendix of the ACSL documentation.

Chapter 20

Misc stuff to move, incorporate or delete

20.0.1 Finding specification files and the refine statement

JML allows specifications to be placed directly in the .java files that contain the implementation of methods and classes. Indeed, specifications such as assert statements or loop invariants are necessarily placed directly in a method body. Other specifications, such as class invariants and method pre- and post-conditions, may be placed in auxiliary files. For classes which are only present as .class files and not as .java files, the auxiliary file is a necessity.

Current JML allows one such auxiliary file per .java file or corresponding .class file. It is similar to the corresponding .java file except that

- it has a .jml suffix
- it contains no method bodies (method declarations are terminated with semicolons, as if they were abstract)
- TBD - field initializations?

The .jml file must be in the same package as the corresponding .java file and has the same name, except for the suffix. It need not be in the same folder, though the tail of the path to the folder containing the .jml file must still correspond to the package containing the .java and .jml files. If there is no source file, then there is a .jml file for each compilation unit that has a specification. All the nested, inner, or top-level classes that are defined in one Java compilation unit will have their specifications in one corresponding .jml file.

The search for specification files is analogous to the way in which .class files are found on the *classpath*, except that the *specspath* is used instead. To find the specifications for a public top-level class *T*:

- look in each element of the *specspath* (cf. section TBD), in order, for a fully-qualified file whose name is *T.jml*. If found, the contents of that file are used as the specifications of *T*.
- if no such *.jml* file is found, look in each element of the *specspath*, in order, for a fully-qualified file whose name is *T.java*.

There are two (silent) consequences of this search algorithm that can be confusing:

- If both a *.jml* and a *.java* file exist on the *specspath* and both contain JML specification text, the specifications in the *.java* file will be (silently) ignored.
- If a *.java* file is listed on the command-line it will be compiled (for its Java content), but if it is not a member of an element of the *specspath*, it will (silently) not be used as the source of specifications for itself.

Obsolete syntax. The *refine* and *refines* statements are no longer recognized. The previous (complicated) method of finding specification files and merging the specifications from multiple files is also no longer implemented. The only specification file suffix allowed is *.jml*; the others — *.spec*, *.refines-java*, *.refines-spec*, *.refines-jml* — are no longer implemented.

In addition, the *.jml* file is now sought before seeking the *.java* file; if a *.jml* file is found anywhere in the *specs path*, then any specifications in the *.java* file are ignored. This is a different search algorithm than was previously used.

20.0.2 Model import statements

This section will be added later. Java `import` statements introduce class names into the namespace of a *.java* file. JML has a `model import` statement:

```
/*@ model import ...
```

The effect of a JML `model import` statement is the same as a Java `import` statement, except that the names imported by the JML statement are only visible within JML annotations. If the `model import` statement is within a *.jml* file, the imported names are visible only within annotations in the *.jml* file, and not outside JML annotations and not in the *.java* file.

Note: Most tools only approximately implement this feature. For example, see `FIXME` for a discussion of this feature in `OpenJML`.

20.0.3 Modifiers

Modifiers are JML keywords that specify JML characteristics of methods, classes, fields, or variables. They are syntactically placed just like Java modifiers, such as `public`.

JML Keyword	Java annotation	class	interface	method	field declaration	variable declaration
code	Code			X		
code_bigint_math	CodeBigintMath					
code_java_math	CodeJavaMath					
code_safe_math	CodeSafeMath					
extract	Extract					
ghost	Ghost				X	X
helper	Helper			X		
instance	Instance					
model	Model					
monitored	Monitored					
non_null	NonNull			X	X	X
non_null_by_default	NonNullByDefault	X	X	X		
nullable	Nullable			X	X	X
nullable_by_default	NullableByDefault	X	X	X		
peer	Peer					
pure	Pure	X	X	X		
query	Query					
readonly	Readonly					
rep	Rep					
secret	Secret					
spec_bigint_math	SpecBigintMath					
spec_java_math	SpecJavaMath					
spec_protected	SpecProtected					
spec_public	SpecPublic					
spec_safe_math	SpecSafeMath					
static	Static					
uninitialized	Uninitialized					

Table 20.1: Summary of JML modifiers. All Java annotations are in the `org.jmlspecs.annotation` package.

[Reuse this table?](#)

Check the table; add section references; add where allowed; indicate which are type modifiers; turn headings 90 degrees.

Obsolete syntax. JML no longer defines the modifier `weakly`.

JML defines some statements that are used in the body of a method's implementation. These are not method specifications per se; rather, they are assertions or assumptions that are used to aid the proof of the specifications themselves, in the way that lemmas are aids to proving a resulting theorem. They can also be used to state predicates that the user believes to be true, and wants checked, or assumptions that are true but are too difficult for the prover to prove itself.

20.0.4 JML expressions

Expressions in JML annotations are Java expressions with three adjustments:

- Expressions with side-effects are not allowed. Specifically, JML excludes
 - the ++ and – pre- and post- increment and decrement operations
 - the assignment operator
 - assignment operators that combine an operation with assignment (e.g., +=)
 - method invocations that are not explicitly declared pure (cf. §TBD)
- JML adds additional operators to the Java set of operators, discussed in subsection §?? below.
- JML adds specific keywords that are used as constants or function-like expressions within JML expressions, discussed in subsection §?? below

JML lock ordering operators (<#) and <#=) The lock ordering operators are used to determine ordering among objects used for locking in a multi-threaded application; the operands are any Java objects. The only predefined property of these operators is that for any two object references o and oo, o <#= oo is equivalent to o == oo || o <# oo; that is <# is like less than and <#= is like less-than-or-equals. There is no predefined ordering among objects. The user must define an intended ordering with some axioms or invariants. An example of using the lock ordering operators for specification and reasoning about concurrency is found in §??.

TBD - add ++ – into the table as Java only; check precedence

20.0.5 Code contracts

This section will be added later.

20.1 JML modifiers and Java annotations

The Java Modeling Language was defined prior to the introduction of annotations in Java. Some, but not all, of the features of JML can now be textually represented as Java annotations. Currently JML supports both the old and new syntactic forms.

Modifiers

Modifiers are JML keywords that specify JML characteristics of methods, classes, fields, or variables. Examples are `pure`, `model`, and `ghost`. They are syntactically placed just like Java modifiers, such as `public`.

<code>new () [] .</code> and method calls		
unary <code>+</code> unary <code>-</code> <code>!</code> (typecast)	-	
<code>*</code> <code>/</code> <code>%</code>	L	
<code>+</code> (binary) <code>-</code> (binary)	L	
<code><</code> <code>></code> <code>>></code>	L	
<code><</code> <code><=</code> <code>></code> <code>>=</code> <code><:</code> <code>instanceof</code> <code><#</code> <code><#=</code>	-	<code><:</code> is the JML subtype operation (§??); <code><#</code> and <code><#=</code> are lock ordering operators (§20.0.4)
<code>==</code> <code>!=</code>	L	
<code>&</code>	L	
<code>^</code>	L	
<code> </code>	L	
<code>&&</code>	L	
<code> </code>	L	
<code>==></code> <code><==</code>	?	JML implies and reverse implies (§??)
<code><==></code> <code><!=></code>	?	JML equivalence and inequivalence (§??)
<code>?:</code>	-	
<code>=</code> <code>*=</code> <code>/=</code> <code>%=</code> <code>+=</code> <code>-=</code> <code><=</code> <code>>=</code> <code>>>=</code> <code>&=</code> <code>⊕=</code> <code> =</code>	L	Java only

Table 20.2: Java and JML operators, in order of precedence, from highest (most tightly binding) to lowest precedence. Operators on the same line have the same precedence. The associativity is given in the central column.

Each such modifier has an equivalent Java annotation. For example

```
/*@ pure */ public int m(int i) ...
```

can be written equivalently as

```
org.jmlspecs.annotation.Pure public int m(int i) ...
```

The `org.jmlspecs.annotation` prefix can be made implicit in the usual way by including the import statement

```
import org.jmlspecs.annotation.Pure;
```

Note that in the second form, the `pure` designation is now part of the *Java* program and so the import of the `org.jmlspecs.annotation` package must also be in the Java program, and the package must be available to the Java compiler.

All of the modifiers, their corresponding Java annotations, and the locations in which they may be used are described in §20.3.5.

Type modifiers

Some modifiers are actually type modifiers. In particular `non_null` and `nullable` are in this category. Thus the description of the previous subsection (§20.1) apply to these as well.

However, Java 1.8 allows Java annotations to be applied to types wherever type names may appear. For example

```
(@NonNull String)toUpper(s)
```

is allowed in Java 1.8 but is forbidden in Java 1.7.

Need additional discussion of the change in JML for Java 1.8, especially for arrays.

Method specification clauses

This section will be added later.

Class specification clauses

This section will be added later.

Statement specifications

JML specifications that are statements within the body of a method have no equivalent as Java annotations. These include loop specifications, assert and assume statements, ghost declarations, set and debug statements, and specifications on individual statements.

20.2 org.jmlspecs.lang package

Some JML features are defined in the `org.jmlspecs.lang` package. The `org.jmlspecs.lang` package is included as a model import by default, just as the `java.lang` package is included by default in a Java file. `org.jmlspecs.lang.*` contains these elements:

- `JML.informal(<string>)` : This method is a replacement for (and is equivalent to) the informal expression syntax `($??) (* ... *)`. Both expressions return a boolean value, which is always `true`.
- TBD

The definition of the Java Modeling Language is contained in the JML reference manual.[?] This document does not repeat that definition in detail. However, the following sections summarize the features of JML, indicate what is and is not implemented in OpenJML, describes any extensions to JML contained in OpenJML, and includes comments about relevant implementation aspects of OpenJML.

20.3 JML Syntax

20.3.1 Syntax of JML specifications

JML specifications may be written as Java annotations. Currently these are only implemented for modifiers (cf. section TBD). In Java 8, the use of Java annotations for JML features will be expanded.

JML specifications may also be written in specially formatted Java comments: a JML specification includes everything between either (a) an opening `/*@` and closing `*/` or (b) an opening `/**@` and the next line ending character (`\n` or `\r`) that is not within a string or character literal.

Such comments that occur within the body of a class or interface definition are considered to be a specification of the class, a field, or a method, depending on the kind of specification clause it is. JML specifications may also occur in the body of a method.

Obsolete syntax. In previous versions of JML, JML specifications could be placed within javadoc comments. Such specifications are no longer standard JML and are not supported by OpenJML.

20.3.2 Conditional JML specifications

JML has a mechanism for conditional specifications, based on a system of keys. A key is an identifier (consisting of ASCII alphanumeric and underscore characters, and beginning with a non-digit). A conditional JML comment is guarded by one or more positive or negative keys (or both). The keys are placed just before the `@` character that is part of the opening sequence of the JML comment (the `/**@` or the `/*@`). Each key is preceded by a `'+'` or a `'-'` sign, to indicate whether it is a positive or negative key, respectively. *No white-space is allowed.* If there is white-space anywhere between the initial `/**` or `/*` and the first `@` character, the comment will appear to be a normal Java comment and will be silently ignored.

The keys are interpreted as follows. Each tool that processes the Java+JML input will have a means (e.g. by command-line options) to specify the set of keys that are enabled.

- If the JML annotation has no keys, the annotation is always processed.
- If there are only positive keys, the annotation is processed only if at least one of the keys is enabled.
- If there are only negative keys, the annotation is processed unless one of the keys is enabled.

- If there are both positive and negative keys, the annotation is processed only if (a) at least one of the positive keys is enabled AND (b) none of the negative keys are enabled.

JML previously defined one conditional annotation: those that began with `/*+@` or `//+@`. ESC/Java2 also defined `/*-@` and `//-@`. Both of these are now deprecated. OpenJML does have an option to enable the `+-style` comments.

The particular keys do not have any defined meaning in the JML reference manual. OpenJML implicitly enables the following keys:

- **ESC** : the ESC key is enabled when OpenJML is performing ESC static checking;
- **RAC** : the RAC key is enabled when OpenJML is performing Runtime-Assertion-Checking.
- **DEBUG** : The DEBUG key is not implicitly enabled. However it is defined as the key that enables the **debug** JML statement. That is the **debug** statement is ignored by default and is used by OpenJML if the user enables the DEBUG key.
- **OPENJML** : The OPENJML key is enabled whenever OpenJML is processing annotations (and presumably is not enabled by other tools).
- **KEY**: The KEY key is reserved for annotations recognized by the KeY tool [?]. It is ignored by OpenJML.

Thus, for example, one can turn off a non-executable assert statement for RAC-processing but retain it for ESC and for type-checking by writing `//-RAC@ assert ...`

20.3.3 Finding specification files and the refine statement

[Discuss obsolete syntax somewhere - including the refines statement](#)

JML allows specifications to be placed directly in the .java files that contain the implementation of methods and `cl3.1sd [(k)10(e)15(y)-348(icussbdeednnot05)-241(to)-241(be)-24(e)such3.1sd [13.1sd [J/F7`

containing the `.java` and `.jml` files. If there is no source file, then there is a `.jml` file for each compilation unit that has a specification. All the nested, inner, or top-level classes that are defined in one Java compilation unit will have their specifications in one corresponding `.jml` file.

The search for specification files is analogous to the way in which `.class` files are found on the *classpath*, except that the *specspath* is used instead. To find the specifications for a public top-level class *T*:

- look in each element of the *specspath* (cf. section TBD), in order, for a fully-qualified file whose name is *T.jml*. If found, the contents of that file are used as the specifications of *T*.
- if no such `.jml` file is found, look in each element of the *specspath*, in order, for a fully-qualified file whose name is *T.java*.

There are two (silent) consequences of this search algorithm that can be confusing:

- If both a `.jml` and a `.java` file exist on the *specspath* and both contain JML specification text, the specifications in the `.java` file will be (silently) ignored.
- If a `.java` file is listed on the command-line it will be compiled (for its Java content), but if it is not a member of an element of the *specspath*, it will (silently) not be used as the source of specifications for itself.

Obsolete syntax. The `refine` and `refines` statements are no longer recognized. The previous (complicated) method of finding specification files and merging the specifications from multiple files is also no longer implemented. The only specification file suffix allowed is `.jml`; the others — `.spec`, `.refines-java`, `.refines-spec`, `.refines-jml` — are no longer implemented.

In addition, the `.jml` file is now sought before seeking the `.java` file; if a `.jml` file is found anywhere in the *specs path*, then any specifications in the `.java` file are ignored. This is a different search algorithm than was previously used.

20.3.4 Model import statements

This section will be added later. Java `import` statements introduce class names into the namespace of a `.java` file. JML has a `model import` statement:

```
//@ model import ...
```

The effect of a JML `model import` statement is the same as a Java `import` statement, except that the names imported by the JML statement are only visible within JML annotations. If the `model import` statement is within a `.jml` file, the imported names are visible only within annotations in the `.jml` file, and not outside JML annotations and not in the `.java` file.

JML Keyword	Java annotation	class	interface	method	field declaration	variable declaration
code	Code			X		
code_bigint_math	CodeBigintMath					
code_java_math	CodeJavaMath					
code_safe_math	CodeSafeMath					
extract	Extract					
ghost	Ghost				X	X
helper	Helper			X		
instance	Instance					
model	Model					
monitored	Monitored					
non_null	NonNull			X	X	X
non_null_by_default	NonNullByDefault	X	X	X		
nullable	Nullable			X	X	X
nullable_by_default	NullableByDefault	X	X	X		
peer	Peer					
pure	Pure	X	X	X		
query	Query					
readonly	Readonly					
rep	Rep					
secret	Secret					
spec_bigint_math	SpecBigintMath					
spec_java_math	SpecJavaMath					
spec_protected	SpecProtected					
spec_public	SpecPublic					
spec_safe_math	SpecSafeMath					
static	Static					
uninitialized	Uninitialized					

Table 20.3: Summary of JML modifiers. All Java annotations are in the `org.jmlspecs.annotation` package.

Note: Most tools only approximately implement this feature. For example, see FIXME for a discussion of this feature in OpenJML.

20.3.5 Modifiers

Modifiers are JML keywords that specify JML characteristics of methods, classes, fields, or variables. They are syntactically placed just like Java modifiers, such as `public`.

Check the table; add section references; add where allowed; indicate which are type modifiers; turn headings 90 degrees.

Obsolete syntax. JML no longer defines the modifier `weakly`.

20.3.6 JML expressions

Expressions in JML annotations are Java expressions with three adjustments:

- Expressions with side-effects are not allowed. Specifically, JML excludes
 - the ++ and – pre- and post- increment and decrement operations
 - the assignment operator
 - assignment operators that combine an operation with assignment (e.g., +=)
 - method invocations that are not explicitly declared pure (cf. §TBD)
- JML adds additional operators to the Java set of operators, discussed in subsection §?? below.
- JML adds specific keywords that are used as constants or function-like expressions within JML expressions, discussed in subsection §?? below

20.3.7 JML types

Specifications are sometimes best written using infinite-precision mathematical types, rather than the fixed bit-width types of Java. JML's arithmetic modes (§??) allow choosing among various numerical precisions. In this section we simply note the type names that JML defines.

All of the Java type names are legal and useful in JML: `int` `short` `long` `byte` `char` `boolean` `double` `real` and `class` and `interface` types. In addition, JML defines the following:

- `\bigint` - the type of infinite-precision integers, represented as `java.lang.BigInteger` during run-time checking
- `\real` - the type of mathematical real numbers, represented as TBD during runtime-checking
- `\TYPE` - the type of JML type objects

The familiar operators are defined on values of the `\bigint` and `\real` types: unary and binary `+` and `-`, `*`, `/`, `%`. Also, these types can be used in quantified expressions and variables of these types can be declared as ghost or model variables.

The set of `\TYPE` values includes non-generic types such as `\type(org.lang.Object)`, fully parameterized generic types, such as `\type(org.utils.List<Integer>)`, and primitive types, such as `\type(int)`. The subtype operator (`<:`) is defined on values of type `\TYPE`.

TBD - what about other constructors or accessors of `\TYPE` values

Chapter 21

Statement translations

TODO: Need to insert both RAC and ESC in all of the following.

21.1 While loop

Java and JML statement:

```
//@ invariant invariant_condition ;
//@ decreases counter ;
while (condition) {
    body
}
```

Translation: *TODO: Needs variant condition, havoc information*

```
{
    //@ assert jmltranslate(invariant_condition) ;
    //@ assert jmltranslate(variant_condition) > 0 ;
    while (true) {
        stats(tmp,condition)
        if (!tmp) {
            //@ assume !tmp;
            break;
        }
        //@ assume tmp;
        stats(body)
    }
}
```

```
    }  
}
```


Chapter 22

Java expression translations

22.1 Implicit or explicit arithmetic conversions

TODO

22.2 Arithmetic expressions

TODO: need arithmetic range assertions

In these, T is the type of the result of the operation. The two operands in binary operations are already assumed to have been converted to a common type according to Java's rules.

```
stats(tmp, - a ) ==>
  stats(tmpa, a )
  T tmp = - tmpa ;
```

```
stats(tmp, a + b ) ==>
  stats(tmpa, a )
  stats(tmpb, b )
  T tmp = tmpa + tmpb ;
```

```
stats(tmp, a - b ) ==>
  stats(tmpa, a )
  stats(tmpb, b )
  T tmp = tmpa - tmpb ;
```

```
stats(tmp, a * b ) ==>
  stats(tmpa, a )
  stats(tmpb, b )
  T tmp = tmpa * tmpb ;
```

```
stats(tmp, a / b ) ==>
  stats(tmpa, a )
  stats(tmpb, b )
  //@ assert tmpb != 0; // No division by zero
  T tmp = tmpa / tmpb ;
```

```
stats(tmp, a % b ) ==>
  stats(tmpa, a )
  stats(tmpb, b )
  //@ assert tmpb != 0; // No division by zero
  T tmp = tmpa % tmpb ;
```

22.3 Bit-shift expressions

TODO

22.4 Relational expressions

No assertions are generated for the relational operations `<` `>` `<=` `>=` `==` `!=`. The operands are presumed to have been converted to a common type according to Java's rules.

```
stats(tmp, a op b ) ==>
  stats(tmpa, a )
  stats(tmpb, b )
  T tmp = tmpa op tmpb ;
```

22.5 Logical expressions

```
stats(tmp, ! a ) ==>
  stats(tmpa, a )
```

T tmp = ! tmpa ;

The `&&` and `||` operations are short-circuit operations in which the second operand is conditionally evaluated. Here `&` and `|` are the (FOL) boolean non-short-circuit conjunction and disjunction.

```
stats(tmp, a && b) ==>
  boolean tmp ;
  stats(tmpa, a)
  if ( tmpa ) {
    /*@ assume tmpa ;
      stats(tmpb, b)
      tmp = tmpa & tmpb ;
    */
  } else {
    /*@ assume ! tmpa ;
      tmp = tmpa ;
    */
  }
```

```
stats(tmp, a || b) ==>
  boolean tmp ;
  stats(tmpa, a)
  if ( ! tmpa ) {
    /*@ assume ! tmpa ;
      stats(tmpb, b)
      tmp = tmpa | tmpb ;
    */
  } else {
    /*@ assume tmpa ;
      tmp = tmpa ;
    */
  }
```

Bibliography

- [1] Gary T. Leavens, Erik Poll, Curtis Clifton, Yoonsik Cheon, Clyde Ruby, David R. Cok, Peter Müller, Joseph Kiniry, Patrice Chalin, and Daniel M. Zimmerman. JML reference manual. Available from <http://www.jmlspecs.org>, September 2009.

Index

(*...*), 53
.jml files, 9
JML.informal(), 53
JML.lbl(), 51
JML.lblneg(), 51
JML.lblpos(), 51
accessible clause, 38
assignable clause, 37
breaks clause, 39
callable clause, 38
captures clause, 39
choose_if clause, 39
choose clause, 39
code, 35
continues clause, 39
diverges clause, 38
duration clause, 38
ensures clause, 37
extract clause, 39
forall clause, 38
ghost, 43
instance, 43
in, 43
maps, 43
measured_by clause, 38
model_program clause, 39
model, 43
monitored, 43
old clause, 38
or clause, 39
peer, 43
query, 43
readonly, 43
rep, 43
requires clause, 37
returns clause, 39
secret, 43
signals_only clause, 37
signals clause, 37
uninitialized, 43
when clause, 38
working_space clause, 38
@Ghost, 43
@Instance, 43
@Model, 43
@Monitored, 43
@Peer, 43
@Query, 43
@Readonly, 43
@Rep, 43
@Secret, 43
@Uninitialized, 43
\TYPE, 27
\bigint, 26
\duration, 58
\elemtype, 55
\exception, 50
\fresh, 53
\invariant_for, 56
\is_initialized, 56
\lblneg, 51
\lblpos, 51
\lbl, 51
\lockset, 58
\max, 58
\nonnulllements, 52
\not_modified, 57
\old, 51
\past, 51
\pre, 51
\reach, 58
\real, 26
\result, 50
\space, 58

- `\typeof`, 55
- `\type`, 54
- `\working_space`, 58
- axiom, 31
- constraint clause, 30
- data groups, 41
- field specifications, 41
- ghost fields, 31
- informal expression, 53
- initializer, 31
- initially clause, 31
- invariant clause, 30
- model classes, 31
- model fields, 31
- model import statement, 28
- model methods, 31
- modifiers, 12
- `non_null`, 41
- `non_null_by_default`, 30
- `nullable`, 41
- `nullable_by_default`, 30
- `package-info.java`, 28
- peer, 42
- represents clause, 31
- `spec_protected`, 41
- `spec_public`, 41
- Specification inheritance, 35
- specification of fields, 41
- `static_initializer`, 31